


RUST CHINA CONE 2023

第三届中国Rust开发者大会



6.17-6.18 @Shanghai

■ Rust原子操作高性能实践

Rust Atomic Deep Dive

Pu Wang @ DatenLord

2023/06/17



Outline

01

What

What are atomic operations in Rust?

02

Why

Why need atomic operations?

03

How

How to understand atomic operations?

04

Memory Model

Memory order in atomic operations

05

Cache Coherence

The overhead of atomic operations

06

Summary

Atomic operation best practice

01

What are Atomic Operations in Rust?

Compare and Swap

```
fn compare_exchange(
    &self, // AtomicI8
    current: i8,
    new: i8,
    success: Ordering,
    failure: Ordering
) -> Result<i8, i8>
```

Fetch and Modify

```
fn fetch_add(&self, val: i8, order: Ordering) -> i8
fn fetch_and(&self, val: bool, order: Ordering) -> bool
...
```

Why Atomic Operation?


High Performance

Lock-free programming

Shared Variable Access When Multi-threading

Lock – context switch

Atomic – no context switch



How to Understand Atomic Operations?



Memory Model

Memory Order

The order of load/store instructions accessing memory.



Cache Coherence

Atomic operation overhead

Atomic operations will change cache line status which might flush cache lines.

Memory Model

- Program order
 - Instructions executed as the order defined in a thread
- Memory order
 - Memory access instructions
 - load & store
 - Out of order
 - Instruction reorder by compilers
 - Out of order execution in CPU
 - load v.s. load
 - load v.s. store
 - store v.s. load
 - store v.s. store

Instruction Reorder by Compilers

```
int A, B;  
  
void foo()  
{  
    A = B + 1;  
    B = 0;  
}
```

```
$ gcc -S -masm=intel foo.c
```

```
$ cat foo.s
```

```
...
```

```
mov     eax, DWORD PTR _B
```

```
add     eax, 1
```

```
mov     DWORD PTR _A, eax
```

```
mov     DWORD PTR _B, 0
```

```
$ gcc -O2 -S -masm=intel foo.c
```

```
$ cat foo.s
```

```
...
```

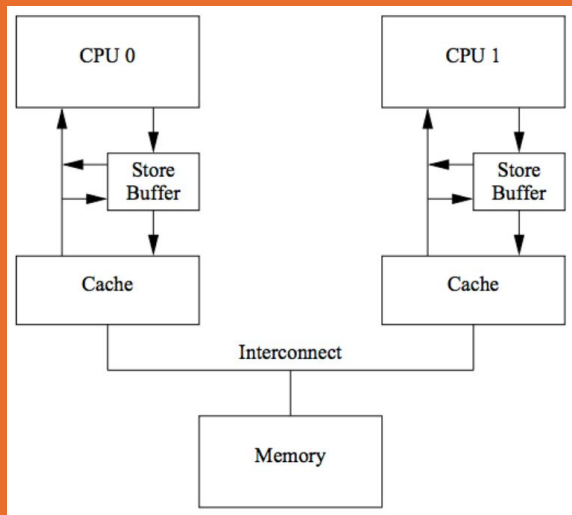
```
mov     eax, DWORD PTR B
```

```
mov     DWORD PTR B, 0
```

```
add     eax, 1
```

```
mov     DWORD PTR A, eax
```


Out of Order Execution



- Write buffer / store buffer
 - Delayed write
 - Store forwarding
- Total store order (TSO)
 - FIFO writer buffer
 - OoO store v.s. load
- Partial store order (PSO)
 - Non-FIFO write buffer
 - OoO store v.s. load
 - OoO store v.s. store

04

Memory Order in C++/Rust

Sequential
Consistency

Acquire

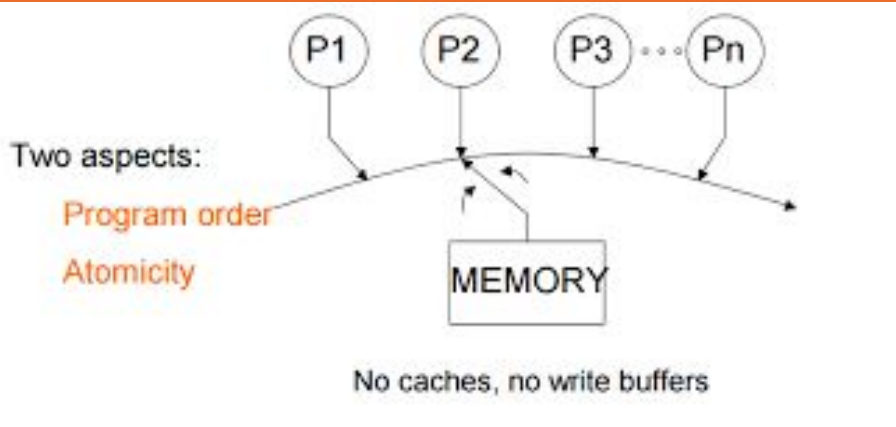
Release

AcqRel

Relaxed

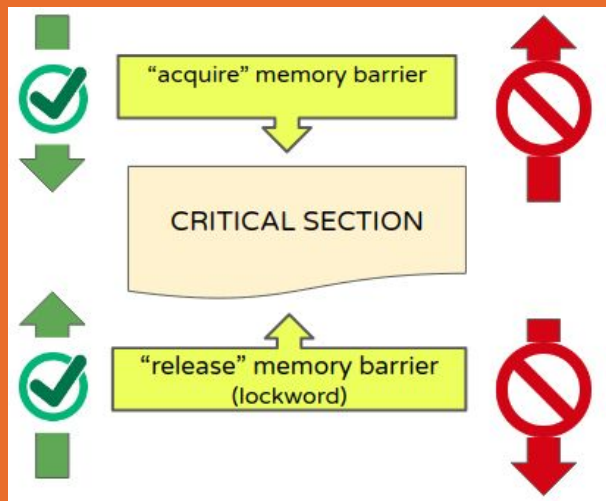
Consume

Sequential Consistency



- Each processor issues memory operations in program order
- The switch provides the global serialization among all memory operations

Acquire Release



- A write-release guarantees that all preceding code completes before the releasing write
- A read-acquire guarantees that all following code starts after the acquiring read

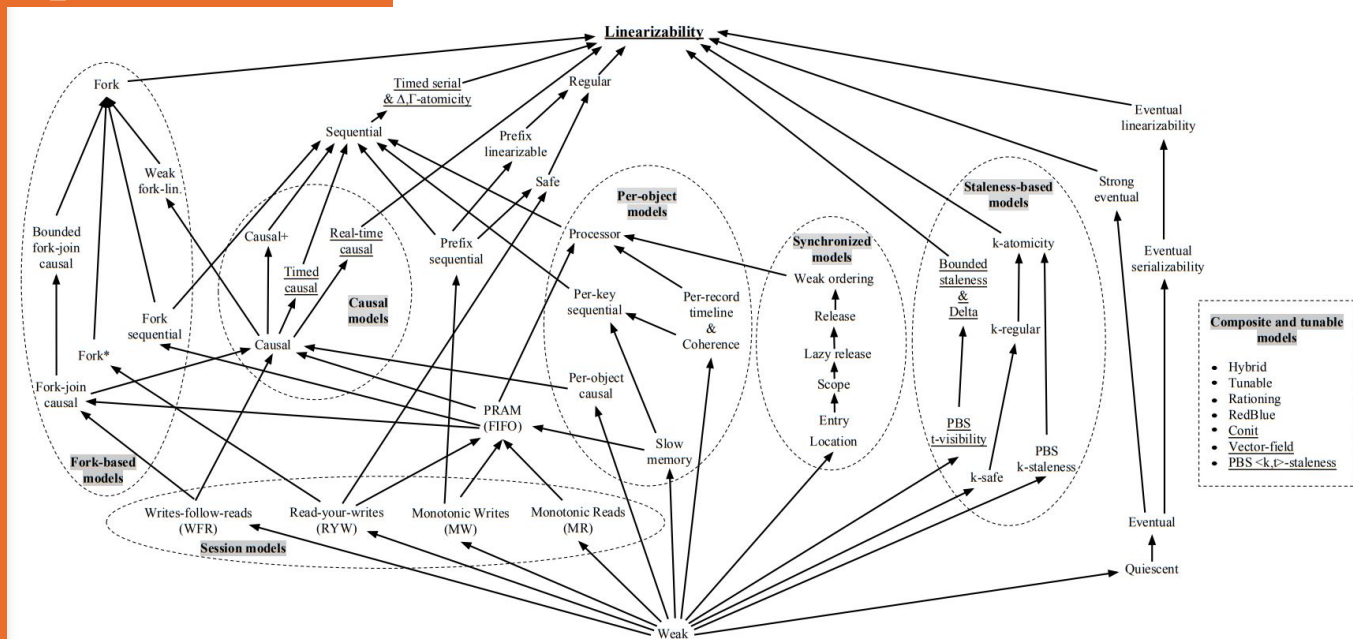
Acquire Release Example

```
std::atomic<std::string*> ptr;  
int data;  
  
void producer()  
{  
    std::string* p = new std::string("Hello");  
    data = 42;  
    ptr.store(p, std::memory_order_release);  
}  
  
void consumer()  
{  
    std::string* p2;  
    while (!(p2 = ptr.load(std::memory_order_acquire)))  
        ;  
    assert(*p2 == "Hello"); // never fires  
    assert(data == 42); // never fires  
}
```

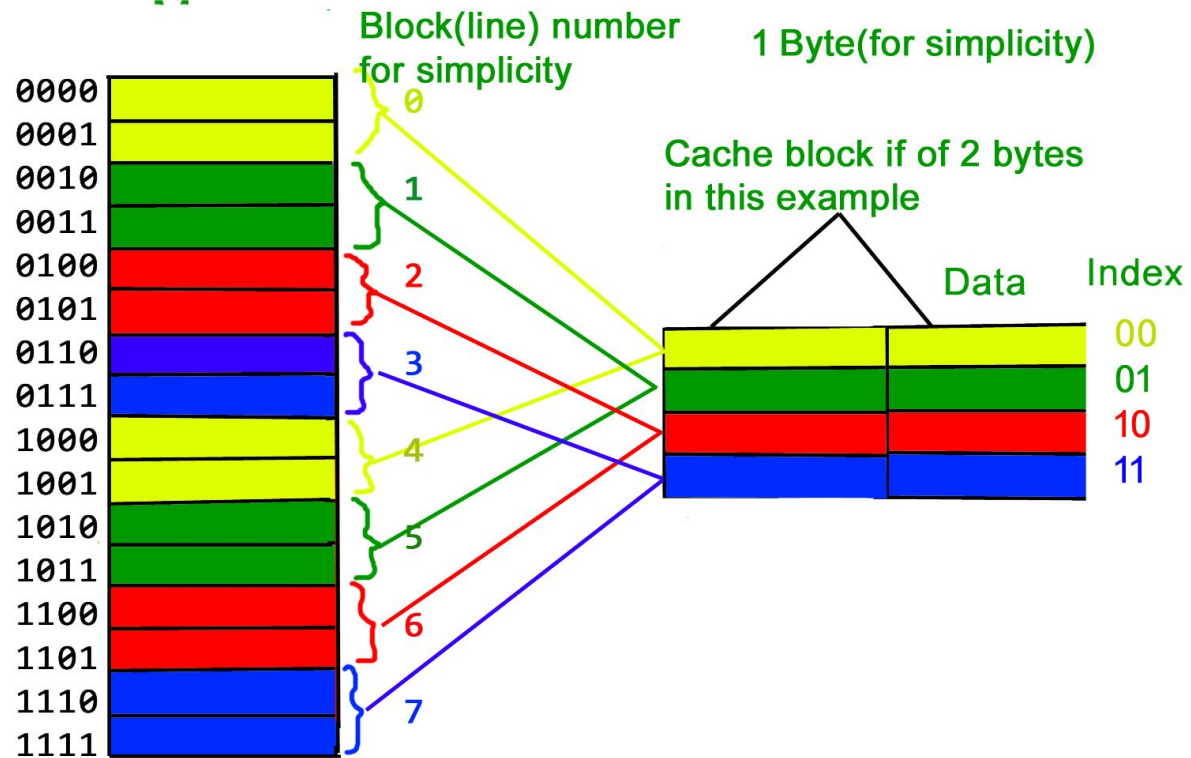
Consistency

Consistency in Non-Transactional
Distributed Storage Systems

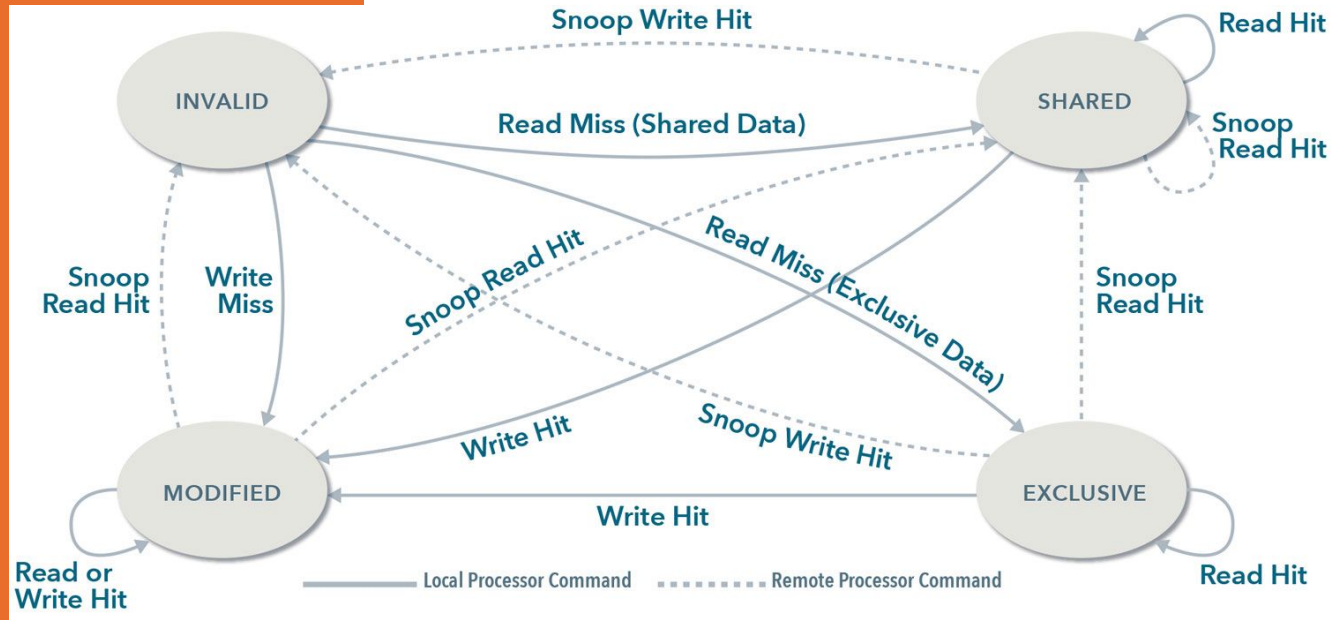
<https://arxiv.org/pdf/1512.00168.pdf>




Cache



Cache Coherence



Exponential Backoff Retry




Summary

Memory Order

SeqCst
Acquire
Release
AcqRel

Cache Coherence

Cache line flush overhead
Exponential backoff retry



Thank you !

